

1. Synthetic Data: How many samples, and types generated?

- **Total number of synthetic datasets:**
TabPFN v2 was pre-trained on approximately 130 million synthetic datasets [arxiv.org+13en.wikipedia.org+13arxiv.org+13](#).
 - **Types of synthetic data:**
The synthetic datasets are generated using a mixture of Structural Causal Models (SCMs) and Bayesian Neural Networks (BNNs). The data-generative prior is designed to reflect:
 - Various causal structures (via SCMs)
 - Functional relationships (BNNs)
 - Feature types: numerical, categorical, missing values
 - Class imbalances and noise distributions [github.comautoml.org+2table-representation-learning.github.io+2en.wikipedia.org+2nature.com+2en.wikipedia.org+2reddit.com+2](#).
-

2. Quantifications: Classes, noise, incompleteness?

Yes—TabPFN explicitly includes:

- **Number of classes**
In v2, datasets may have up to 10 classes [arxiv.org+8nature.com+8paperswithcode.com+8arxiv.org+1paperswithcode.com+1](#).
 - **Noise**
 - Random noise is added at all nodes in the causal networks (SCMs) and to outputs in BNNs [table-representation-learning.github.io+1github.com+1](#).
 - Noise distributions are sampled per synthetic dataset.
 - **Missing/incomplete data**
 - TabPFN's prior incorporates missing values and imbalanced targets, mirroring real-world tabular data [superdatascience.com+11en.wikipedia.org+11table-representation-learning.github.io+11](#).
 - **Feature counts & scales**
 - Training covers up to 500 features and up to 10,000 samples during evaluation [nature.com+1reddit.com+1](#).
-

3. How is synthetic data fed into the model—offline or on-the-fly?

- **TabPFN uses offline pre-training:**
 - The 130 million synthetic datasets are entirely generated before training.
 - Training proceeds by sampling these datasets and optimizing the transformer to approximate Bayesian posterior predictive distributions on held-out test splits (in-context learning) [arxiv.org+14table-representation-learning.github.io+14automl.org+14arxiv.org+5en.wikipedia.org+5reddit.com+5](#).

- During inference, no additional synthetic data is generated or training steps taken. Instead, TabPFN performs a single forward pass with your dataset (training + test examples as input tokens) .

Summary Table

Aspect	Details
Number of datasets	~130 million synthetic datasets generated (offline)
Data types	From SCMs & BNNs, includes numerical, categorical, missing, imbalanced
Classes	Up to 10 classes
Noise	Injected at each causal/Bayesian node
Missing/incomplete	Modelled synthetically, training includes missing data
Feeding method	Generated offline → used in pre-training; no on-the-fly generation
Inference	Single forward pass via in-context learning

Why these matters

TabPFN's strength comes from learning a generic, Bayesian-inspired learning algorithm thanks to exposure to a massive and diverse synthetic data universe. This gives it robustness to real-world dataset idiosyncrasies—without needing dataset-specific tuning.

Here's what I found:

Is the synthetic data generation code publicly available?

While the TabPFN GitHub repository and related projects are publicly accessible, the actual code used to generate the ~130 million synthetic training datasets is not included. As noted by the authors:

“The code to run inference is available, just not the training code and data generation.”

(news.ycombinator.com)

In other words, the model and inference tools are open-source, but the training data script—the generator used to create the synthetic SCM/BNN datasets—is not released.

Implications for your analysis

Since the generation code is not available, we cannot inspect exact implementation details or run tools like Gemini 2.5 on that code—there is nothing to ingest or analyze for verification. Therefore:

- My prior answers (based on the paper and supplementary) remain accurate.
- No code-level clarifications exist (e.g., precise distributions, parameter settings) beyond what is described in the paper.

- Any deeper confirmation of assumptions or edge-case behavior inside the synthetic data generator is impossible without access to the code.
-

✅ Summary

- The public TabPFN code includes model inference, but not synthetic-data generation.
- Training generation scripts are closed—likely proprietary.
- Without that code, we can't reconfirm details via inspection or external tools.

Let me know if you'd like guidance on possible workarounds, or deeper dives based on the paper's supplementary!

No, TabPFN does *not* create synthetic data during training on your small or medium dataset.

🔍 Here's the full explanation:

✅ Pretraining Phase (Before You Use It)

- TabPFN is pretrained on 130 million synthetic datasets.
- These are generated once, using Structural Causal Models (SCMs) and Bayesian Neural Networks (BNNs).
- This training teaches the model *how to learn from small datasets*.

🔧 This happens offline — you don't see this or need to do anything.

🚫 What *doesn't* happen

When you give TabPFN your small or medium-sized dataset:

- It does not generate synthetic data.
 - It does not perform any fine-tuning.
 - It does not do any retraining or gradient descent.
-

🧠 What it *does* do instead

- TabPFN treats your dataset (few labeled + unlabeled samples) as a sequence.
- It passes this sequence through the transformer in one forward pass.
- It predicts the labels for the unlabeled (test) examples.

This is called in-context learning, and it's very fast.

📌 Summary Table

Phase	Does it generate synthetic data?	Purpose
Pretraining (offline)	 Yes	Teach TabPFN how to learn
Inference (your dataset)	 No	Predict using what it learned